

CONTACT CENTER · MEMBER SERVICES AND PHARMACY

Agent Assist for advocates

A working prototype, and the thinking behind it.

Harish Krishnan

Humana Take Home Prototype · 09.21.2026

Why I built Agent Assist



The riskiest place in the business

The live call is the highest-volume, most regulated thing Humana does with a member.



Five jobs, from memory

The exact legal disclaimers, facts from six systems, the right flow, objections, documentation. All while talking to someone.



Found weeks later, if ever

A miss turns up in a small QA sample, long after the member hung up.



A chatbot can't fix it

A legal requirement isn't a question anyone asks.
It has to follow the call and speak first.

SO IT EXISTS TO DO ONE THING

Find the mistake while the member is still on the phone, not weeks later.

Who Agent Assist is for

Week one

Still in training.
Someone is sitting
with them

OUR USER

About six months in

Runs the call alone. Hasn't memorised the
exceptions yet.

Five years in

Knows it cold. Gains
mainly from the call
notes being drafted
for them



They feel it most. Five things at
once, one of them a word-for-word
legal disclaimer



It's where the miss happens. Fixing
it in QA afterwards fixes the report,
not the call



The group is growing. Advocates
leave sooner than they used to, so
there are more new ones every year

Not the primary user: QA and team leads, who read the call afterwards. The same approved disclaimers could score their calls later.

Advocate journey · Today vs Agent Assist

Stage	Today	With Agent Assist	What the member gets
Before call	The phone menu says why the member called. Often that is not the real reason	Opens the moment the call connects, greeting ready	—
During call	Six systems to search. A new question buries the last. A paraphrased disclaimer goes unnoticed	One card, every fact sourced. Missing words shown if one slips. Open questions stay open	Less time on hold while someone hunts for an answer
After call	Notes typed from memory. The outcome code guessed	Notes drafted from what happened, each line naming its source	Not asked the same things again by the next team

THE PATTERN

Everything the advocate was holding in their head is now on the screen, with its source attached.

The business opportunity

FIRST



Compliance

A missed disclaimer caught and corrected on the call. Every call, not a sample.

THEN



Capacity

No hunting through systems mid-call. Notes drafted in seconds.

THEN



Consistency

Same question, same answer, same source, whoever picks up.

THEN



Getting up to speed

New advocates reaching full speed sooner.

SIZING THE OPPORTUNITY · THREE SUMS, AND THE COST OF FAILURE IS THE FIRST

What a failure costs

Disclaimers a year × today's miss rate × the cost of one failure

Time back per call

After-call minutes saved, plus shorter calls × calls a year × the cost of an advocate hour

Training a new advocate

Weeks to full speed × advocates hired a year × the cost of a week

WHAT THE PILOT WILL MEASURE

How often a disclaimer gets missed. How many minutes come back per call. How many weeks a new advocate needs. **Those three numbers, in dollars and in time.**

What Agent Assist is measured on

NORTH STAR METRIC

Disclaimer
miss rate

=

Legal disclaimers the advocate
never said

Legal disclaimers that were required

Counted **per disclaimer, not per call**. One call owing three counts three times.

👍 Success KPIs · time and effort saved

After-call work time · the clearest win

Systems opened by hand, per call

Time to first fact · question to answer on screen

First Contact Resolution · no call-back needed

Time to proficiency · weeks to full speed

⚠️ Guardrails · must not get worse

Handle time · watched, never a target

Member experience · survey score and complaints

Advocate trust · suggestions used, not skipped

Enrolment rate · read beside declined offers

AI actions in regulated steps: zero

Make shorter calls the target and advocates rush the disclaimer, which is the exact failure this product exists to catch. That is why handle time is watched and never a goal.

Operational metrics · the full set

🕒 Should go down

After-Call Work · minutes spent on notes after the member hangs up

Application switches · systems opened by hand to answer one question

Time to first fact · seconds from the question to the answer on screen

Hold time · minutes the member spends on hold

Time to proficiency · weeks before a new advocate runs calls like an experienced one

👉 Should go up

First Contact Resolution · sorted without the member calling back

Transfer accuracy · transfers that reach the right team first time, so nobody re-explains

Answer consistency · same question, same answer, whoever picks up

Agent adoption · advocates using it, and how often they take a suggestion

⚠️ Watched, never a goal

Average Handle Time · must not get worse, and nobody is given a goal to cut it

Enrolment rate · reported beside complaints and declined offers

Captured from · the call log, the phone system and contact records

WHY HANDLE TIME IS NOT A GOAL

Make shorter calls the target and advocates rush the disclaimer. That is the exact failure this product exists to catch.

Who captures value



Compliance

More required disclaimers actually said, and evidence on every call instead of a sample.

BECOMES MONEY WHEN

Legal puts a figure on what one finding costs.

The member gains too, and no budget line holds it. It surfaces later as fewer call-backs.



Operations

Minutes back after every call, and no hunting through systems during one.

BECOMES MONEY WHEN

Those minutes are rebuilt into the staffing plan, not absorbed as slack.

Not claiming a savings percentage or a headcount number. Neither has a starting figure.



The wider business

One approved list of legal wording, one set of playbooks. The next channel uses them instead of writing its own.

BECOMES MONEY WHEN

A second channel ships without rebuilding them.

It also costs: model spend per call, and someone owning bad suggestions and the wording.

Agent Assist · in scope vs not in scope

✓ Built now

- Transcription (synthetic call transcripts)
- Legal disclaimers detected and prompted
- Knowledge retrieval and query routing
- Answers with source and citations
- Suggestions, transfer, outcome code
- Human-only enrolment
- End-of-call notes
- Evaluation and test framework

🔑 Next

- Phone platform and live speech to text
- Humana's member, pharmacy and benefits records
- Member authentication
- Monitoring, security and logging
- Cost and speed measured per call
- Testing and evaluation on production calls
- Post-call quality scoring for QA
- **Responsible AI review and Legal sign-off**
- A small launch with senior advocates first

🔒 Never

- The AI signing anyone up
- The AI deciding coverage
- Taking payment
- Clinical judgement
- Writing legal disclaimers
- Ranking advocates

Whoever asks, and however good the model gets

THE END STATE

Every regulated call, in every advocate's desktop. **One approved list of legal wording**, owned by Legal and operations. Watched at every step, and **scored against real calls**.

Agent Assist today · the thinking and the AI are real. **Data and integration are mocked**.

Prototype

What you just watched was real, and what was not

⚡ Real AI · changes with what is said

- Reading what the caller means
- Spotting when a disclaimer is due
- Choosing which records to open
- Searching the documents
- Drafting answers, notes and suggestions
- Checking each sentence against its record

☁ Faked · and labelled as such

- The phone system and speech to text
- Checking who the caller is
- Benefits, claims, pharmacy and provider records
- The approved wording library
- The transfer actually connecting

🔒 Neither · it is plain code

- The legal wording, copied word for word
- When each one is due, and what counts as said
- The words on screen when one is missed
- The word checks that spot a likely disclaimer
- The seven things that block a suggestion
- The three things that block an outcome code

Known limitations: no after-call quality scoring · it can run slower than target under load · this is not production security, throughput or integration.

Agent Assist architecture

[link to the diagram goes here]

Legal disclaimers

1



Legal owns the words

A small official list: three exact sentences, plus rules for when each is required, how often it must be said, and what counts as a miss.

The sentences are locked. Nobody rewrites them on the fly.



2



Code decides when

Agent Assist never writes a disclaimer, and never decides the official text.

It notices when a conversation reaches a moment a rule covers: a future price being quoted, an optional pharmacy service actually being discussed. Fixed word checks catch the obvious cases. The model is asked only when they can't tell.



3



Code checks what was said

The exact sentence appears on screen. It stays there until the advocate reads it.

Code compares what the advocate said against that sentence and marks it **said, late, paraphrased or still missing**.

A past price, an amount the member says, or words on screen the advocate never read out do not count.

Grounded answers · where the facts come from

1



The model picks what to look up

It chooses which records to open.
Nothing is scripted in advance.



2



Every sentence names its record

One sentence, **one record** it came from.



3



Code re-reads the record

Code opens it again and checks the amounts, dates and names.



4



Or the sentence never appears

If the record does not carry them, **the advocate never sees it.**

Next Best Action

1



The menu is fixed

- Approved steps only
- The model cannot invent one
- A playbook says when each fits

At most one. Often none.



2



The model drafts

- Reads this call's facts
- Writes words to say
- Never a new action
- No price before consent

Playbook says none, it stays none.



3



Fixed rules sit in front

- Identity not checked
- A disclaimer is due
- Already enrolled, or said no
- Asked not to be contacted
- Built on an unconfirmed fact

Nothing shows. Code decides, not the model.



4



The advocate decides

- Sign someone up
- Decide coverage
- Take payment
- Give medical advice

The AI has no tool for any of these. Not blocked. Not there.

Offer, ignore or confirm. A no ends it.

Important tradeoffs

Type	What we chose	What we gained	What we lost
Design	A due disclaimer takes the whole panel	The advocate cannot miss it under something more interesting	Additional cues are queued, one click away. Some offers never get made
Design	One card at a time, never a dashboard	While talking, the advocate reads one thing	Detail is one click away, not already in view
Architecture	Knowledge retrieval: the AI picks which systems to check	It handles questions nobody planned for in advance	Slower. Cached records, shorter answers, faster model: 6.5 to 1.8 seconds
Architecture	Every sentence must name its source	Every fact traces to a record. Invented amounts are stripped out	Checking is almost free. Writing the source cost 1.8 seconds
Architecture	Fetch a record, read it, then pick the next	Much simpler, and all ten two-step questions came back right	Needs rethinking if answers fail because two records do not connect



Obligations loud, optional help quiet. The shaded row is the one the other four follow from.

How Agent Assist is judged

The job	What counts as right	Who checks it
Showing the required disclaimer	The exact words, before they are due, once per call	Plain code. No false alarms
Reading what the caller says	Tells a clear yes from a maybe, and never hides a real question	Code, against lines a person has already marked up
Answering, suggesting, writing the notes	Facts right, each one sourced, and it says so when there is no answer it can support	Code for the facts. A second AI for the wording
Finding the right passage in a document	The right passage near the top, the wrong plan rejected	Code, on questions we already know the answer to

Where an AI does the grading, a person spot-checks the grader. **If they disagree on more than one answer in ten, we don't trust the grades.**

What was tested

Run over and over while the logic was built. **These have been run.**

The whole Harry call

Every beat from greeting to outcome code, including the missed disclaimer caught and read late

Legal disclaimers

Right speaker, once per call, past charges do not trigger one, unclear words neither pass nor fail

The member says no

The suggestion goes and does not come back. Nothing tries to talk them round. The disclaimers still get read

A record is missing

It shows what the member paid, says the cause is not confirmed, and guesses at nothing

The pharmacy changes mid-lookup

The estimate already in flight is thrown away, never relabelled

Made-up numbers

An invented amount is stripped out, and a real record that does not hold the fact is caught

The AI tries to sign someone up

It cannot. Only the advocate can, and the permission is one-time and tied to exactly what was read out

Someone trying to break in

No member data before identity is checked. An instruction hidden in a record changes nothing

Nothing is canned

Same question, different member, different answer. Adding a new offer means adding a document, not code

What the testing taught. Automated tests caught broken logic. They never caught what the screen showed. **That second job cannot be handed to the agent that wrote the code.**

The evaluation set

36

Answers

All five members. Their own records, the documents, both, two-step questions, and questions the documents don't answer

19

Suggestions

The right one or none. Each of the seven things that block a suggestion. Quiet after a no

13

Must never do

Clinical advice, AI enrolment, another member's data, pressure. **One failure blocks release**

30

Said differently

Ten questions, three phrasings each. Same facts either way

11

Change the data

Alter a record and the answer changes with it. Nothing is canned

Roadmap · the path forward

Phase	What happens	Before we move on
Measure today 4 weeks	Run the evaluation on calls Humana has already recorded. Show the screen to newer advocates on replays. No live call is touched	Our scoring agrees with Humana's own reviewers, in every safety category
Hardening and integration 8 weeks	Phone, live text from audio, member records read-only, monitoring, and a carry-on-without-it mode	Speed and cost per call within budget. Security and privacy sign off
Quiet run 4 weeks	The whole thing on real calls, showing nobody anything. We score what it would have said	The checker matches the reviewers on real audio, with few false alarms
Beta program 6 weeks	Ten to twenty experienced advocates who can tell when it's wrong, and will say so. Wording and rules get corrected from what they report	Accuracy targets met, no safety incidents, and they keep using it unprompted
The people it's for 8 weeks	One team of newer advocates, alongside a similar team without it	Fewer misses left uncaught, call length no worse, steady use. Operations and Compliance both approve

WHY EXPERIENCED ADVOCATES GO FIRST

They catch the mistakes. They know the right answer, so a wrong one gets reported instead of read out.

They decide if the floor trusts it. Advocates believe the person beside them, not a rollout email.

One of them ends up owning it. The wording needs a champion inside operations, not a vendor ticket.

A

Appendix

Twelve slides. Every number in this deck has its source one slide away.

- Testing assumptions
- The five deliverables
- The four must-shows
- Evaluation dimensions
- The price comparison
- The three required disclaimers
- Architecture detail
- Taking Agent Assist live
- Measurements
- How the KPI is counted
- The optional stretch
- Decision register

APPENDIX

Testing assumptions

About	We assumed	How we'd find out we were wrong
Members	Harry is a character we wrote, not a real pattern in your calls	Listen to real calls. He stops being the design case
Advocates	Bad suggestions early cost more trust than good ones later win back	Watch what they skip. If untrue, show suggestions sooner
The desktop	There's room for this beside what they already keep open	Ask the desktop team. A narrow strip shrinks the layout
Legal	Said late but correctly beats never said	Ask Legal. If they're the same, the headline changes

What I'd do differently

Pick the call type **before** writing the spec. Build something running sooner. Fix the screen before the logic.

Still open

Several model calls run at the same time. I'd hunt for timing bugs before anyone relied on this.

APPENDIX A1

The five deliverables

Brief asks for	The slide	What it establishes
1 · Problem Brief	"Why I built Agent Assist" and "Who Agent Assist is for"	Five simultaneous jobs, why a question-answering panel isn't enough, why now, and the group the design is for
1b · its business opportunity	"The business opportunity" and "Who captures value"	Compliance, capacity, consistency and ramp time, the three sizing sums with the missing numbers named, and whose budget each benefit lands in
2 · Service Design	"Advocate journey · today vs Agent Assist"	Before, during and after, each with the intended outcome and what the human still owns throughout
3 · Working Prototype	"Prototype", then "What you just watched was real, and what was not", then the four how-it-works slides	All four must-shows, running live against simulated systems with invented records, with the simulated parts named before anyone asks
4 · Product Roadmap	"Roadmap", with "Taking Agent Assist live" in the appendix	Five phases, each with the question it answers and a named approver before moving on, and why experienced advocates go first
5 · Key Tradeoff	"Important tradeoffs", the first row	Obligations loud, optional help quiet, with the cost named and the five smaller choices shown as supporting it

"Show your thinking" isn't collected at the end. Decisions sit on "Important tradeoffs" and "In scope vs not in scope". Assumptions and what I'd do differently have their own slide, "Testing assumptions", at the front of this appendix. Past experience is placed where the lesson landed, in the notes for "Grounded answers", rather than gathered into a career slide.

The four must-shows

Brief requirement	How it is built	Where you see it
a · Live call context	Timed stream carrying partial, final, corrected and uncertain events. Call strip shows identity, call type and the current need with its flow step. Call type changes from evidence, refill into education and enrollment, never by beat number	Beats 1–5
b · Correct verbatim in flow	Greeting, pricing and closing text pulled byte for byte from the simulated scripting registry. While a disclaimer is due, the Now card is that wording only. No quote amounts on that card	Beats 1, 10, 14
c · Real-time nudge on a miss	Two-stage trigger, then a code string match. A paraphrase shows what was said, what is missing, what was added, visible to the advocate without prototype details. Late exact delivery is retained as a finding	Beat 10
d · Auto wrap and handoff	Drafted from session evidence, each sentence naming one source and support-checked. Disposition recommended from a governed document with cited reasons. Connection is not coverage resolution	Beat 14

Also on the same call, from the brief's "great looks like": intent and call-type recognition from evidence · grounding in the member's own records · a governed objection reply the advocate chooses to use · a warm transfer with a stated role reason · scoped human-only enrollment.

Evaluation dimensions

Dimension	Slides that carry it	What does the scoring
Problem and user insight	Why I built Agent Assist · Who Agent Assist is for · The business opportunity	Simultaneous workload, why a chatbot can't fix it, the group and why it grows
AI that does real work	Legal disclaimers · Grounded answers · Next Best Action · the prototype	A two-step trigger, the model choosing its own lookups, a source checked for every statement, and suggestions that must pass code first
Journey mapping	Advocate journey, today vs Agent Assist · Roadmap	Before, during and after with intended outcomes, then the path to live
Execution	Prototype · What you just watched was real, and what was not · How Agent Assist is judged · What was tested · The evaluation set	A running call, the simulated parts named before anyone asks, four jobs each checked a different way, and regression suites that have run kept apart from a set that has not
Judgment	Important tradeoffs · In scope vs not in scope · Roadmap	One lead tradeoff with its cost, three scope boundaries, phase order with named approvers
Responsible AI	Next Best Action · Legal disclaimers · In scope vs not in scope	Wording never judged by a model, blocks enforced in code per kind of action, and four things the AI has no tool for at all
Thinking and accumulated judgment	Speaker notes throughout · Testing assumptions	Experience placed where each lesson landed. Assumptions get their own slide, each paired with what would disprove it, plus what I'd do differently
Architecture and platform fit not in the brief's table; my own row	Agent Assist architecture · Next Best Action · Taking Agent Assist live	Content kept apart from the logic, one doorway to the data, and one approved disclaimer list feeding the call, the scorecard and other channels

APPENDIX A5

The price comparison

Future 90-day supply	Lakeview standard retail	Oak Street preferred retail	CenterWell delivery
Atorvastatin 20 mg · 90 tablets	\$15	\$6	\$6
Metformin 500 mg · 180 tablets	\$60	\$24	\$18

WHAT THE ARITHMETIC MEANS

On atorvastatin, preferred retail and delivery are **equal**. CenterWell is not uniquely cheapest.

On metformin, CenterWell is **\$6 below preferred retail**, not \$42 below every retail option. \$42 is only the gap to Lakeview.

A selective comparison framed as the complete cheapest choice is a hard failure, not a quality issue. Omitting the relevant retail option and claiming best would fail the reference contract.

These amounts are **independently supplied estimates** for that exact drug, quantity, supply period, plan and pharmacy. They are not derived by multiplying the historical charges, and they guarantee nothing.

No amount appears on any card **until comparison consent is a clear yes**. Asked before consent, "why was it \$27?" mentions no future price.

The three required disclaimers

Requirement	Exact disclaimer	Applicability and deadline	Repetition and recovery
Greeting	"Thank you for calling Humana. This call is being recorded."	At the opening, before identity questions or substantive servicing	Once per interaction. Correct delivery is quiet, no acknowledgement click
Pricing	"Any price estimate we discuss is based on the information available today and may change when your prescription is filled."	Before the advocate first speaks a prospective estimate. A dollar amount, or comparative price language between named options. Excludes completed charges and a member mentioning prices	One adequate reading covers later estimates, both drugs and corrected quotes. A topic switch does not reset it
Closing v2	"Your decision today has no impact on your plan membership."	Once the advocate actually introduces or discusses the optional service choice, including answering a member-initiated inquiry. Due before exit or handoff	A clear reread when the attempt is uncertain. Attestation is defined but deferred, not built

An internal suggestion, a call-type label, or an Offer click does not trigger the closing.
The actual discussion does. Obligations come from conversation evidence, never from a click.

Reference outcome for this call: 2 of 3 fully satisfied, 1 pricing-timing finding recovered on the call, 0 unable to verify.
The finding is not erased by the correction.

Each record carries the label **"Synthetic prototype policy – not actual Humana-approved wording."** The model never writes, edits or judges this text.

Architecture detail

THE RETRIEVAL BOUNDARY

Six simulated systems over HTTP: **benefits, eligibility, claims, pharmacy, provider, scripting**. Telephony and coverage review sit **beside** them, not inside claims.

App code **never imports fixture files**. Records load only inside the simulated route handlers, so the retrieval path is real even though the data is not.

Every response carries its source system, an as-of and a simulated flag. Evidence labels read **System record · pharmacy · simulated**.

TWO TIERS, AND WHAT EACH DOES

Fast tier: per-utterance interpretation, and disclosure trigger classification on ambiguous partials. Warmed at call connect.

Mid tier: the answer tool loop, ready-answer generation, wrap, handoff, next best action, disposition.

Embeddings: documents embedded once at startup, query embedded per search. **No vector database.**

DELIBERATELY ABSENT

No graph. A 10-question chain evaluation completed 10/10 under eight seconds with no failure from a missing hop, so the decision was recorded rather than assumed.

No mining pipeline, no dashboard, no vector store, no agents framework. Simplest evidence-complete path.

No pre-written answers. Ready answers are generated from the documents and keyed to a source hash, so they cannot go stale silently.

Evaluator isolation is enforced, not documented. Golden answers live only under the test folder, and a lint rule plus a build alias stop application code from importing them. Not a comment-only rule. Payload shapes are illustrative: real integration maps to Humana's APIs, and nothing here is NCPDP, CMS or FHIR compliant.

APPENDIX

Taking Agent Assist live

Hearing the call

Real audio, turned into text as it happens, with the speakers told apart

Who's calling

Humana's own identity check, handed over by the phone system

Member records

The real systems. Reading first, writing notes later. Permissions, and a log of who saw what

The other systems

Pharmacy prices, the customer record, the quality tool, one login for all of it

Watching it run

Alerts when it breaks, cost and speed at each step, health details hidden in the logs

The documents

Owners, versions, approval. Legal owns the required disclaimers and their rules

The models

Hosting Humana approves for health data, and rules on what gets kept

Where it lives

Inside the advocate's own desktop, opening when the call connects

Safety

People trying to break it on purpose, and an "it's down, carry on" mode

Proof it works

Scoring built from recorded calls and checked against Humana's own reviewers

People and sign-off

Training, telling advocates plainly what is watched, and four teams signing off

What a prototype proves

That this is worth building. Not that it survives a Tuesday. Roughly my first two quarters

APPENDIX A10

Measurements

Measurement	Number	What it actually measures, and how to cite it
Governed utterance rules	14–26 ms	A conservative pattern match on the transcript line, timed on the 90-day and comparison-yes rules. No model call. This is the 1-second-safe code path, and the pricing stage-1 trigger runs on it
Trigger classification	median 878 ms, max 1,176 ms, 2 of 20 over 1s	Ambiguous partials only, on the priority tier, hedged with two requests at 2× list price. Enum output, not interpretation
Per-utterance interpretation	23 of 23 finals over 1s. Median 1,327 ms, max 3,967 ms	Reading call type, needs, focus and consent from real utterances on the accepted run. This is the number for that path
Keep-alive ping	median 899 ms, 4 of 10 over 1s	A bare 16-token ping to the model endpoint. Not a product path. Do not cite it as a warm median
Chain evaluation	10/10 correct, 0 over 8s	10 multi-step questions across five members on one dated run. This is the run behind the no-graph decision
Chain, later runs	3 further runs, 10/10 each. One question at 9,647 ms	The same suite under load. A real miss, kept. Diagnosed as load, not a missing join, so the graph decision stands
Speed levers	6,549 ms → 1,778 ms median	Flagship answer. Load the snapshot once, keep answers short, faster tier. Measured one lever at a time
Support check	0.1–3 ms per statement	Re-reading the cited record and comparing values. Not a cause of a timing miss by itself
Model calls in one window	588 calls, 588 succeeded, mean 1,709 ms	One measured test window. Not an availability claim

APPENDIX A11

How the KPI is counted

THE UNIT

One **required disclaimer**, not a call, an alert, a reread or a model decision. A single call can owe three disclaimers and each is counted separately.

Reported **by requirement type**, so a greeting problem cannot hide inside a pricing average.

THREE OUTCOMES PER INSTANCE

Satisfied on time · **Failed** with evidence · **Unable to verify**. The three add up to every known-applicable instance.

Failures split into **recovered on the call** and **unrecovered**. A paraphrase alone is unrecovered.

A closing cannot be recovered once the advocate exits. It is unrecovered by definition.

HOW IT COULD BE GAMED

Unrecovered falls while **unable to verify rises**. That is not better compliance, it is a smaller denominator.

Recovered-late failures rise while unrecovered falls. That is **advocates leaning on the nudge**, and it is a regression.

The sampled call mix, the disclaimer definitions or the denominator change **after** results are seen. Fix all three first.

THREE THINGS I SAY OUT LOUD

1 · The prototype demonstrates the instruments, not the result. One fictional call cannot produce this number. A pilot gives the first real figure.

2 · The headline assumes Legal treats a disclaimer read late but exactly as better than one never read. If they treat those the same, the headline becomes total failures, and the value rests on detection speed and on every call being checked rather than sampled.

3 · There is no baseline miss rate and I will not quote one. The first deliverable of a pilot is to establish it, by running the checker silently on calls before advocates see anything.

The optional stretch

HOW IT WOULD WORK

The QA view shares the **same approved disclaimers, applicability rules and version** that guided the advocate live.

It reconstructs from the underlying transcript, speaker and timing evidence, and system actions. **Not from the live badges** or the generated wrap.

Three outcomes per instance: fully satisfied, finding evidenced, or unable to verify. Each finding annotated recovered or unrecovered.

THE THREE CHECKS THAT WOULD PROVE IT

Reconstruct the call from evidence and obtain the reference counts.

Tamper: turn every live badge green and delete the pricing obligation object, leaving the raw evidence alone. QA must still find the deadline failure.

Version: supply a newer rule version. The earlier call must not be silently regraded under it.

WHY IT IS NOT BUILT

It is optional and comes **after** the required prototype and the core checks. Omitting it does not make the core incomplete.

In two to three days it would have come out of the deferred answer, the nudge, or the scoped enrollment, and those carry the argument.

The other stretch option, a ramp-time estimate, I did not model. I will not supply a number I cannot support.

WHAT IT WOULD NOT PROVE

Recomputing from the same governed source does **not** give statistically independent errors. A shared bad source stays a risk, and human review is not removed. It is a review aid, never employee ranking, discipline or certification.

Decision register

#	Decision	Downside, or what could change it
1	Compliance-led whole-call assistance. No prevention claim	Several outcomes need separate measurement
2	Visible call type and flow, plus independent open needs	More state than a linear script
3	One main call, six short contrasts, compact boundary checks	The sample is intentionally dense
4	Exact text and governed use together; timing and uncertainty assessed separately	Recovery and review burden. Real approved policy may differ
5	Embedded five-region workspace, one current card, direct recovery	Some useful context is less prominent
6	Answer, clarify, partial help, withhold or stay silent, by evidence	Partial answers can look complete
7	Separate safety, usefulness, timing and effort; bounded tests	A small synthetic set can overfit
8	Optional per-call QA reconstructed from evidence	Shared-source errors and review cost persist
9	Lead with reviewed unrecovered failures; total failures as the guardrail	Depends on Legal treating a late reading as better than none
10	Pilot a real compliance opportunity, deepen, then scale	Narrow scope delays broader value
11	Obligations loud, optional help quiet — the named tradeoff	Useful optional assistance may be missed
12	Two experience lessons with explicit claim limits	Less prototyping emphasis than a career summary would give
13	One fixed fictional world: the charges, the comparison, the election, the pending case	More fixture detail, and no claim it is real operations
14	Validate the work, the rules, the speech, the users and the outcomes	Findings can overturn the prototype design. That is a reason to test