


Humana Agent Assist

A copilot for the advocate on a member call. It follows the conversation, answers what is asked, tracks the wording the advocate owes, and proposes at most one next action.

1 Next.js app · 8 simulated services · 19 governed documents · 3 required disclosures · 7 code gates

The whole system



Inside the brain

Green serves the member. Yellow watches the call.

Knowledge retrieval

RAG

- **19 documents.** Playbooks, articles, policies, an FAQ. They sit behind a service, not in the prompt.
- **Common questions are answered ahead of time,** saved with a vector at start-up. A later question matches by meaning in one step.
- **The rest go the long way.** Code drops documents that do not fit the plan, then the model reads what is left.
- **Nothing ships unquoted.** A sentence survives only if a source backs it. If none survive, the screen says so.

Required disclosures

OBLIGATIONS

- **Three sentences the advocate owes the member.** The recorded-call notice, the pricing disclaimer, and a line saying a no costs them nothing.
- **Each has its own due time.** The notice before identity questions, the pricing line before the first price.
- **The rail tracks them live.** One chip each: read on time, read now, said late, could not hear it, missed.
- **The check is word by word.** Code first, a model second, and the log says which.

Next best action

NBA

- **Ten playbooks, filtered to the plan.** A compliance change is a document edit.
- **The model picks from the list,** or picks none. An invented action is rejected.
- **Facts must already be confirmed.** A card citing something the call never established is never drawn.
- **Seven checks in code can kill it,** then a person presses the button. The AI never enrolls anyone, approves coverage, takes payment, or gives medical advice.
- **The seven:** identity not verified, required wording due, already enrolled, member said no, do-not-contact flag, already dismissed, cites an unconfirmed fact.

Quality score

AFTER-CALL QA

- **Rebuilt on every read.** Touches no live state, writes nothing, decides nothing.
- **Five counts.** Required, on time, late and fixed, never fixed, could not tell.
- **Each miss shows its clock.** When it fell due, how fast the nudge came, how long to the right words.

Data layer

THE THREE ABOVE RUN ON THIS

- **One way in.** Every line of the call lands at a single endpoint, whether the script played it or the presenter typed it.
- **One way out.** Eight simulated Humana services. Nothing here opens a file. It asks a service and gets data stamped with the system that answered and the time it was read.
- **Held for the call.** After identity is verified the member is fetched once, in parallel: plan, cost share, prescriptions, claims, preferences. Refill status is the exception and is always fetched fresh.
- **Held across calls.** The document answers and their vectors, built at start-up and saved to one file keyed by a hash of the source documents. Edit a document and the set rebuilds itself.
- **Locked to the caller.** Every lookup uses the session's member id. The model cannot ask about anyone else.
- **Written down as it happens.** Lookups, stops and timings go straight to the log beside it.


Log data

THE SCORE RUNS ON THIS

- **One file per call.** Every lookup, block, timing and disclosure status change.
- **No totals are stored,** so the score and the log cannot drift apart.
- **Click a number** to open the lines behind it.

One question, end to end

Where the time goes when a member asks something, and what the advocate sees while it goes there. The stand-in services answer in milliseconds, so the shape of this path is real but its clock is not. Against live systems, the parallel start and the cached lane are what would keep the wait bearable.



The slow parts start before the model asks for them, and the advocate gets words to say while they run. A cached hit answers in one vector comparison; anything member-specific goes the long way, round by round, with several calls fired in parallel each time.

The models

Two language models and an embedding model, each picked for its job. All three are set by environment variable, so any of them can be swapped.

OpenAI Luna

FAST

`gpt-5.6-luna`

- **Reads every line as it arrives.** What the member is asking for, whether a yes is a real yes, whether the advocate just said something that starts a clock.
- **Chosen for speed.** It sits in the path of the call, so its cost is paid on every utterance.
- **Never the last word.** Code checks first and Luna is the second look. The log records which of the two decided.

OpenAI Terra

HEAVIER

`gpt-5.6-terra`

- **Does the work that has to be right.** Chooses and chains the lookups, writes the answer, drafts the next-best-action card.
- **Also runs at start-up,** writing the ready answers the fast lane matches against later.
- **Fenced in.** It picks tools, never members. Actions from a list, never invented ones. Its sentences are deleted if nothing retrieved backs them.

OpenAI embeddings

VECTORS

`text-embedding-3-small`

- **Turns text into numbers so meaning can be compared.** Questions and saved answers both become vectors.
- **This is what makes the fast lane fast.** One comparison instead of a model call and a tool loop.